LLARS: Enabling Domain Experts & Developers Collaboration for LLM Prompting, Generation and Evaluation

Philipp Steigerwald[1], **Mara Stieler**[2], **Jennifer Burghardt**[2], **Eric Rudolph**[1], **Jens Albrecht**[1]
Technische Hochschule Nürnberg Georg Simon Ohm, Nürnberg, Germany
[1]Faculty of Computer Science, Center for Artificial Intelligence (KIZ)
[2]Faculty of Social Sciences, Institute for E-Counselling
{philipp.steigerwald, mara.stieler, jennifer.burghardt, eric.rudolph, jens.albrecht}@th-nuernberg.de

Abstract

We demonstrate LLARS, an open-source platform that bridges the gap between domain experts and developers for building LLM-based assistants. LLARS integrates three tightly connected modules into an end-to-end pipeline: **Collaborative Prompt Engineering** for real-time co-authoring with version control and instant LLM testing, **Batch Generation** for configurable output production across user-selected prompts × models × data with cost control and **Evaluation** for human and LLM evaluators to rate, rank, label or compare outputs with live inter-rater agreement metrics. New prompts and models are automatically available for batch generation and completed batches can be turned into evaluation scenarios with a single click. Interviews with six domain experts and three developers in online counselling confirmed that LLARS eliminates the "translation work" between disciplines and enables domain experts to directly co-author prompts and conduct evaluations independently.

Demo: https://llars-project.github.io/demo

1 Introduction

Organizations building LLM-based assistants for sensitive domains [Chen *et al.*, 2024; Ling *et al.*, 2025]—online counselling [Stade *et al.*, 2024], legal analysis [Katz *et al.*, 2024], medical documentation [Van Veen *et al.*, 2024]—must write effective prompts, select suitable models and rigorously evaluate outputs, requiring both domain expertise and technical skill [Zamfirescu-Pereira *et al.*, 2023; Schulhoff *et al.*, 2024]. Public benchmarks rarely transfer to domain-specific tasks; practitioners must test models on their own data to find the smallest—and most cost-effective—option that still meets quality requirements, especially when replacing a model in a running system. Yet domain experts and developers typically work in isolation: prompts are iterated in shared documents without version control and no single tool covers the full cycle from prompt development through systematic evaluation.

We present LLARS (LLM Assisted Research System), an open-source platform that closes this gap with three tightly connected modules (Figure 1). *Collaborative Prompt Engineering* lets domain experts and developers co-author prompts in a shared real-time editor with version control and instant LLM testing. *Batch Generation* produces the user-configured Cartesian product of prompts × models × data items with cost estimation and budget control. *Evaluation* lets human and LLM evaluators rate, rank, label or compare outputs with live agreement metrics. New prompts and models automatically appear in batch generation and completed batches become evaluation scenarios with a single click; each module also supports standalone JSON/CSV export and connects to any OpenAI-compatible provider.

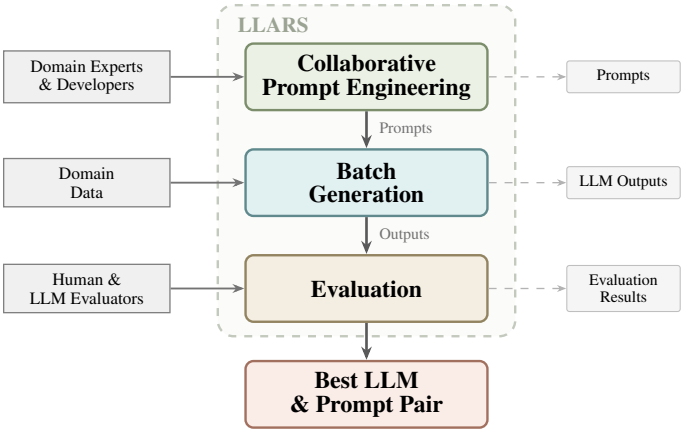

Figure 1: LLARS pipeline: domain experts and developers collaboratively develop prompts, generate outputs across LLMs and evaluate results with human and LLM evaluators. Each stage supports export; the pipeline yields a validated model–prompt combination.

Table 1 surveys twelve existing platforms along these dimensions. Agenta [Agenta AI, 2024] combines prompt versioning with LLM-as-judge scoring and provides A/B deployment testing, but lacks real-time co-editing and structured human evaluation with inter-rater agreement metrics. Phoenix [Arize AI, 2024], Langfuse [Langfuse GmbH, 2024] and W&B Weave [Weights & Biases, 2024] provide observability dashboards with dataset experiments and per-output feedback but focus on tracing and debugging production systems rather than systematic batch generation or coordinated multi-evaluator campaigns. ChainForge [Arawjo *et al.*, 2024]

Table 1: Comparison of GUI-first tools for prompt engineering, batch generation and hybrid LLM output evaluation. ● = native support; ◐ = partial; ○ = not supported. **Collab.** = real-time collaborative prompt co-editing; **Batch** = systematic generation over prompts × models × data; **H-Eval** = structured human evaluation campaigns; **LLM-E** = LLM-as-evaluator; **IRR** = inter-rater reliability metrics; **E2E** = integrated end-to-end pipeline.

Tool	OSS	Collab.	Batch	H-Eval	LLM-E	IRR	E2E
LLARS	●	●	●	●	●	●	●
Agenta [Agenta AI, 2024]	●	◐	◐	◐	●	○	◐
ChainForge [Arawjo *et al.*, 2024]	●	○	●	○	◐	○	◐
Phoenix [Arize AI, 2024]	●	○	◐	◐	●	○	○
Langfuse [Langfuse GmbH, 2024]	●	○	◐	◐	●	◐	○
W&B Weave [Weights & Biases, 2024]	●	○	◐	◐	●	○	○
Label Studio [HumanSignal, Inc., 2024]	●	○	○	●	◐	◐	○
Argilla [Argilla, Inc., 2024]	●	○	○	●	◐	◐	○
LangSmith [LangChain, Inc., 2024]	○	○	◐	◐	●	○	◐
Braintrust [Braintrust Data, Inc., 2024]	○	○	◐	◐	●	○	◐
Maxim [Maxim AI, 2024]	○	○	◐	◐	●	○	◐
Vellum [Vellum AI, 2024]	○	○	◐	◐	●	○	◐

offers a visual node-based interface for multi-model comparison with template variables and supports programmatic evaluation via custom code nodes, but omits structured human evaluation entirely. Label Studio [HumanSignal, Inc., 2024] and Argilla [Argilla, Inc., 2024] offer dedicated human annotation workflows—Label Studio computes Krippendorff's α across annotators, Argilla provides agreement analytics via its API—but both require externally generated outputs, disconnecting evaluation from prompt development. LangSmith [LangChain, Inc., 2024], Braintrust [Braintrust Data, Inc., 2024], Maxim [Maxim AI, 2024] and Vellum [Vellum AI, 2024] provide playground-style testing with dataset-level scoring but are closed-source, lack collaborative editing and omit batch generation. LLARS is an open-source platform to unify collaborative prompt engineering, Cartesian-product batch generation and hybrid human–LLM evaluation with agreement analytics, enabling developers and domain experts to work together from prompt development through rigorous assessment.

2 Platform Overview

Co-designed with developers, social science researchers and practising counsellors, LLARS runs as a containerized web application for textual data. Teams register proprietary or self-hosted OpenAI-compatible LLM providers; registered models then appear across all three modules.

2.1 Collaborative Prompt Engineering

The collaborative editor (Figure 2), built on Yjs [Jahns, 2023], lets every keystroke appear instantly for all connected users. A prompt is composed of ordered blocks—one optionally designated as system prompt, the rest concatenated into the user prompt—each maintaining its own version history that tracks insertions and deletions per save (visible as `+130 / −0` in Figure 2), enabling diff comparison and rollback without affecting other blocks.

Blocks may contain template variables (`{{variable_name}}`) as placeholders for external

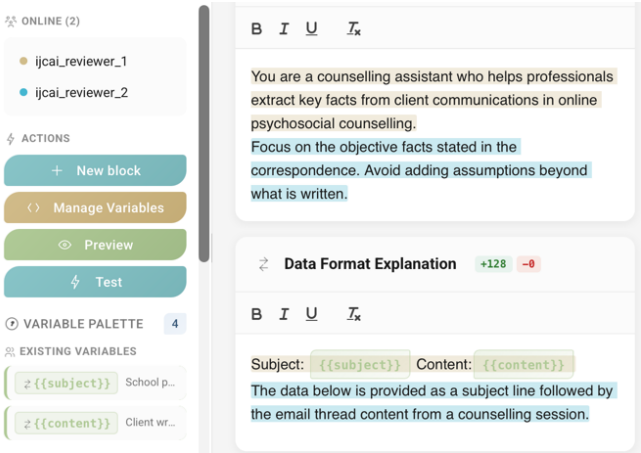

Figure 2: Collaborative prompt editor with a data-format explanation block containing template variables `{{subject}}` and `{{content}}`. The sidebar (left) shows actions for block management, variable management, preview and LLM testing. The Variable Palette (bottom-left) lists all extracted variables with sample values that can be dragged into blocks.

data—an entire email thread, for instance, can be captured in a single `{{content}}` variable, keeping the surrounding instructions concise. All variables are collected in the shared Variable Palette (Figure 2, bottom-left) and annotated with sample values (e.g. `{{subject}}` = *"School problems at home"*) that are substituted when testing against a selected LLM, with responses streamed in real time so that multiple team members can test different models in parallel. Prompts are exportable as JSON and directly available in batch generation, where template variables are filled from dataset columns to scale a single prompt to hundreds of outputs.

2.2 Batch Generation

While the prompt editor supports rapid single-example testing, a single example cannot validate whether a prompt gen-

eralizes across diverse inputs. Batch generation addresses this: users select from available prompts—self-authored or shared by collaborators—and any registered model to apply to their data, producing the Cartesian product of prompts $\times$ models $\times$ data items. A generation matrix previews all resulting combinations and estimates cost based on model pricing and projected token counts; optional budget caps pause jobs that exceed a threshold. A wizard offers three data-source modes: upload (CSV/JSON columns map to the prompt's template variables), previous generation (reuse data from a completed batch to test new prompts or models on the same dataset) and prompt only (generate without source data, filling variables from sample values). For example, 50 counselling email threads combined with two prompts and two models produce $50 \times 2 \times 2 = 200$ outputs.

Results stream in real time and every output records its full provenance—source item, prompt version, model, parameters, token count and cost—exportable as CSV or JSON. Evaluators see only the generated text without any indication of which model or prompt produced it; completed batches transfer into evaluation with a single click, preserving all attribution for downstream provenance analysis.

2.3 Evaluation

Each evaluation is organized as a *scenario*—a self-contained campaign that defines the type, assigns evaluators and distributes items—supporting six types: rating, ranking, labeling, comparison, mail rating and authenticity. Scenarios are created from batch outputs with a single click or via an AI-assisted Scenario Wizard; owners control whether all items go to every evaluator or configurable subsets are distributed for faster throughput.

Where most tools limit LLMs to pairwise judging [Zheng *et al.*, 2024], LLARS treats LLM evaluators as full participants alongside humans: they receive the same items, use the same interface and their judgments enter the same agreement analysis.

The default evaluation mode is *bucket ranking*: evaluators sort items into ordinal categories (e.g. *Good*, *Medium*, *Bad*) and rank within each bucket—reducing cognitive load [Miller, 1956] and yielding more reliable judgments than total ordering [Kiritchenko and Mohammad, 2017], as confirmed by domain expert interviews. Traditional ranking remains available.

Once assessments are submitted, scenario owners inspect aggregated results live, including automatically computed inter-rater reliability (e.g. Krippendorff's α). The central outcome is the *provenance analysis* (Figure 3): since every item carries its generating model and prompt from batch generation, LLARS ranks each combination by top-bucket hit rate—directly answering *which pair should we deploy?*

3 Deployment and User Study

LLARS is actively used in online counselling research: the Virtual Client project [Rudolph *et al.*, 2024] uses it for prompt development of simulated client interactions, and CAIA [Steigerwald *et al.*, 2025], an AI assistant for email counselling, relies on it for prompt development and evaluation. A full end-to-end case study [Steigerwald and Albrecht, 2025] produced 253 subject lines across 11 LLMs and had five professionals with an LLM evaluator rate them (1,518 assessments), determining how small a model could be while still meeting quality requirements.

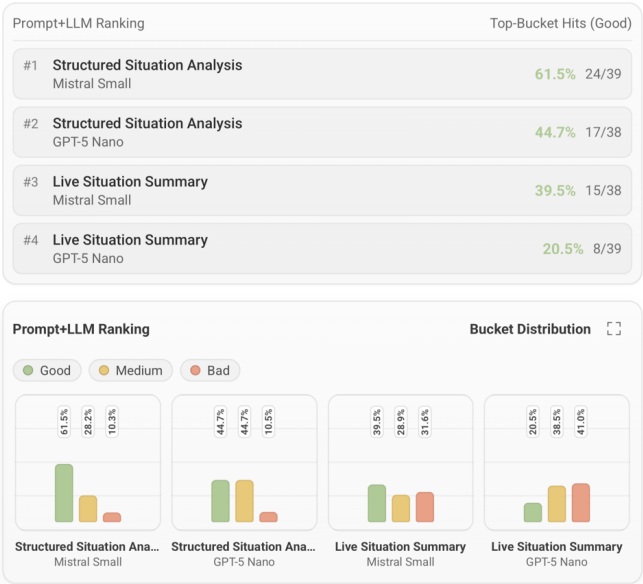

Figure 3: Provenance analysis ranking model–prompt combinations by top-bucket hit rate (top) with per-combination bucket distributions (bottom).

Semi-structured interviews with six domain experts and three developers confirmed that consolidating the pipeline in one platform saves considerable time. Domain experts found prompt authoring straightforward and evaluation immediately understandable; developers valued that prompts, data and outputs remain in one place. Both groups emphasised that co-locating the workflow removed the "translation work" between disciplines.

4 Demo and Conclusion

Our demo walks attendees through the complete workflow—from collaborative prompt authoring through batch generation to evaluation with live agreement metrics and provenance analysis—using pre-built counselling scenarios that let attendees author prompts, trigger batch runs and evaluate outputs hands-on.

By unifying prompt engineering, batch generation and evaluation in a single open-source platform, LLARS closes the loop from prompt development to rigorous assessment. While our current deployments focus on online counselling, the platform is domain-agnostic and applicable to any text-based LLM evaluation task. Unlike existing tools that address individual stages (Table 1), LLARS lets prompts flow into generation, outputs into evaluation, and provenance traces results back to their source—all within a shared workspace accessible to domain experts and developers alike. Interviews confirmed that this integration eliminates the "translation work" between disciplines and enables domain experts to contribute directly at every stage.

References

[Agenta AI, 2024] Agenta AI. Agenta: Open-source LLMOps platform for prompt management, evaluation, and observability. https://agenta.ai, 2024. Accessed: 2026-02-01.

[Arawjo *et al.*, 2024] Ian Arawjo, Chelse Swoopes, Priyan Vaithilingam, Martin Wattenberg, and Elena Glassman. ChainForge: A visual toolkit for prompt engineering and LLM hypothesis testing. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. ACM, 2024.

[Argilla, Inc., 2024] Argilla, Inc. Argilla: Collaboration tool for AI engineers and domain experts. https://argilla.io, 2024. Accessed: 2026-01-06.

[Arize AI, 2024] Arize AI. Phoenix: Open-source AI observability and evaluation. https://phoenix.arize.com, 2024. Accessed: 2026-02-01.

[Braintrust Data, Inc., 2024] Braintrust Data, Inc. Braintrust: AI evaluation and observability platform. https://www.braintrust.dev, 2024. Accessed: 2026-02-01.

[Chen *et al.*, 2024] Zhiyu Zoey Chen, Jing Ma, Xinlu Zhang, Nan Hao, An Yan, Armineh Nourbakhsh, Xianjun Yang, Julian McAuley, Linda Petzold, and William Yang Wang. A survey on large language models for critical societal domains: Finance, healthcare, and law. *Transactions on Machine Learning Research*, 2024.

[HumanSignal, Inc., 2024] HumanSignal, Inc. Label studio: Open source data labeling platform. https://labelstud.io, 2024. Accessed: 2026-01-06.

[Jahns, 2023] Kevin Jahns. Yjs: A CRDT framework for shared editing. https://yjs.dev, 2023. Accessed: 2026-01-06.

[Katz *et al.*, 2024] Daniel Martin Katz, Michael James Bommarito, Shang Gao, and Pablo Arredondo. GPT-4 passes the bar exam. *Philosophical Transactions of the Royal Society A*, 382(2270):20230254, 2024.

[Kiritchenko and Mohammad, 2017] Svetlana Kiritchenko and Saif Mohammad. Best-worst scaling more reliable than rating scales: A case study on sentiment intensity annotation. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 465–470. Association for Computational Linguistics, 2017.

[LangChain, Inc., 2024] LangChain, Inc. LangSmith: AI agent evaluation platform. https://www.langchain.com/langsmith, 2024. Accessed: 2026-02-01.

[Langfuse GmbH, 2024] Langfuse GmbH. Langfuse: Open-source LLM engineering platform. https://langfuse.com, 2024. Accessed: 2026-02-01.

[Ling *et al.*, 2025] Chen Ling, Xujiang Zhao, Jiaying Lu, Chengyuan Deng, Can Zheng, Junxiang Wang, Tanmoy Chowdhury, Yun Li, Hejie Cui, Xuchao Zhang, et al. Domain specialization as the key to make large language models disruptive: A comprehensive survey. *ACM Computing Surveys*, 58(3), 2025.

[Maxim AI, 2024] Maxim AI. Maxim: GenAI evaluation and observability platform. https://www.getmaxim.ai, 2024. Accessed: 2026-02-01.

[Miller, 1956] George A. Miller. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological Review*, 63(2):81–97, 1956.

[Rudolph *et al.*, 2024] Eric Rudolph, Natalie Engert, and Jens Albrecht. An AI-based virtual client for educational role-playing in the training of online counselors. In *Proceedings of the 16th International Conference on Computer Supported Education (CSEDU)*, pages 108–117, 2024.

[Schulhoff *et al.*, 2024] Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, HyoJung Han, Sevien Schulhoff, et al. The prompt report: A systematic survey of prompting techniques. *arXiv preprint arXiv:2406.06608*, 2024.

[Stade *et al.*, 2024] Elizabeth C. Stade, Shannon Wiltsey Stirman, Lyle H. Ungar, Cody L. Boland, H. Andrew Schwartz, David B. Yaden, Joao Sedoc, Robert J. DeRubeis, Robb Willer, and Johannes C. Eichstaedt. Large language models could change the future of behavioral healthcare: A proposal for responsible development and evaluation. *npj Mental Health Research*, 3:12, 2024.

[Steigerwald and Albrecht, 2025] Philipp Steigerwald and Jens Albrecht. Comparing large language models for automated subject line generation in e-mental health: A performance study. In *Proceedings of the 11th International Conference on Information and Communication Technologies for Ageing Well and e-Health (ICT4AWE)*, pages 70–77. SciTePress, 2025.

[Steigerwald *et al.*, 2025] Philipp Steigerwald, Nico Bienlein, Jennifer Burghardt, Mara Stieler, Robert Lehmann, and Jens Albrecht. CAIA in practice: Field evaluation of an AI-assisted support system for text-based online counselling. In *Proceedings of the IEEE International Conference on Tools with Artificial Intelligence (ICTAI)*, 2025. Accepted.

[Van Veen *et al.*, 2024] Dave Van Veen, Cara Van Uden, Louis Blankemeier, Jean-Benoit Delbrouck, Asad Aali, Christian Bluethgen, Anuj Pareek, Malgorzata Polacin, Eduardo Pontes Reis, Anna Seehofnerova, et al. Adapted large language models can outperform medical experts in clinical text summarization. *Nature Medicine*, 30:1134–1142, 2024.

[Vellum AI, 2024] Vellum AI. Vellum: AI application development platform. https://www.vellum.ai, 2024. Accessed: 2026-02-01.

[Weights & Biases, 2024] Weights & Biases. Weave: Toolkit for developing AI-powered applications. https://wandb.ai/site/weave, 2024. Accessed: 2026-02-01.

[Zamfirescu-Pereira *et al.*, 2023] J.D. Zamfirescu-Pereira, Richmond Y. Wong, Bjoern Hartmann, and Qian Yang.

Why johnny can't prompt: How non-AI experts try (and fail) to design LLM prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, 2023.

[Zheng *et al.*, 2024] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36, 2024.