LLARS: Enabling Domain Experts & Developers Collaboration for LLM Prompting, Generation, and Evaluation

Philipp Steigerwald[1], **Mara Stieler**[2], **Jennifer Burghardt**[2], **Eric Rudolph**[1], **Jens Albrecht**[1]
Technische Hochschule Nürnberg Georg Simon Ohm, Nürnberg, Germany
[1]Faculty of Computer Science, Center for Artificial Intelligence (KIZ)
[2]Faculty of Social Sciences, Institute for E-Counselling
{philipp.steigerwald, mara.stieler, jennifer.burghardt, eric.rudolph, jens.albrecht}@th-nuernberg.de

Abstract

We demonstrate LLARS, an open-source platform that enables domain experts and developers to collaboratively build, test, and evaluate LLM-based systems. Developing such systems for sensitive domains requires both domain knowledge—understanding what constitutes quality—and technical skill for prompt design and model configuration, yet these groups typically lack a shared environment. LLARS addresses this with three integrated modules: **Collaborative Prompt Engineering** lets both groups co-author prompts in a real-time editor with version control and instant LLM testing. **Batch Generation** systematically produces the Cartesian product of prompts × models × data items with cost tracking. **Evaluation** supports human and LLM evaluators assessing outputs via rating, ranking, labeling, or comparison with live inter-rater agreement metrics. Each module operates independently or as part of an end-to-end pipeline. LLARS is actively used for developing AI-assisted systems in psychosocial counselling, where domain experts and developers jointly engineer prompts, compare models, and evaluate outputs.

Demo: https://llars-project.github.io/demo

1 Introduction

Organizations building LLM-based assistants for sensitive domains [Chen *et al.*, 2024; Ling *et al.*, 2025]—psychosocial counseling [Stade *et al.*, 2024], legal analysis [Katz *et al.*, 2024], medical documentation [Van Veen *et al.*, 2024]—must write effective prompts, select suitable models, and rigorously evaluate outputs, requiring both domain expertise (understanding *what* constitutes quality) and technical skill (structuring prompts and configuring LLM pipelines) [Zamfirescu-Pereira *et al.*, 2023; Schulhoff *et al.*, 2024]. Yet these two groups typically work in isolation: prompts are iterated in shared documents without version control, and no single tool supports the full cycle from collaborative prompt development through systematic evaluation.

We present LLARS (LLM Assisted Research System), an open-source co-designed platform that closes this gap with three integrated modules (Figure 1). *Collaborative Prompt Engineering* lets domain experts and developers co-author prompts in a shared real-time editor with version control and instant LLM testing. *Batch Generation* systematically produces the Cartesian product of prompts × models × data items with cost estimation and budget control. *Evaluation* lets human and LLM evaluators assess outputs side by side—via rating, ranking, labeling, or pairwise comparison—with live agreement metrics (Krippendorff's α, Cohen's κ, and others). Each module can be used independently with JSON/CSV export, and the platform connects to any OpenAI-compatible provider including locally hosted models.

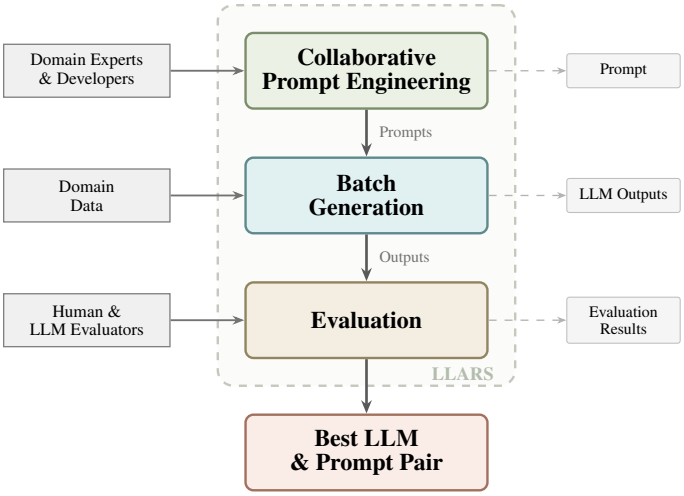

Figure 1: LLARS pipeline: domain experts and developers collaboratively develop prompts, generate outputs across LLMs, and evaluate results with human and LLM evaluators. Each stage supports export; the pipeline yields a validated model–prompt combination.

Table 1 surveys twelve existing platforms along these dimensions. Agenta [Agenta AI, 2024] combines prompt management with LLM-as-judge scoring but lacks real-time co-editing and structured human evaluation with agreement metrics. Phoenix [Arize AI, 2024], Langfuse [Langfuse GmbH, 2024], and W&B Weave [Weights & Biases, 2024] provide dataset experiments and per-output feedback but focus on tracing rather than systematic batch generation or multi-evaluator campaigns. ChainForge [Arawjo

Table 1: Comparison of GUI-first tools for prompt engineering, batch generation, and hybrid LLM output evaluation. ● = native support; ◐ = partial; ○ = not supported. **Collab.** = real-time collaborative prompt co-editing; **Batch** = systematic generation over prompts × models × data; **H-Eval** = structured human evaluation campaigns; **LLM-J** = LLM-as-judge; **IRR** = inter-rater reliability metrics; **E2E** = integrated end-to-end pipeline.

Tool	OSS	Collab.	Batch	H-Eval	LLM-J	IRR	E2E
LLARS	●	●	●	●	●	●	●
Agenta [Agenta AI, 2024]	●	◐	◐	◐	●	○	◐
ChainForge [Arawjo *et al.*, 2024]	●	○	●	○	◐	○	◐
Phoenix [Arize AI, 2024]	●	○	◐	◐	●	○	○
Langfuse [Langfuse GmbH, 2024]	●	○	◐	◐	●	○	○
W&B Weave [Weights & Biases, 2024]	●	○	◐	◐	●	○	○
Label Studio [HumanSignal, Inc., 2024]	●	○	○	●	◐	◐	○
Argilla [Argilla, Inc., 2024]	●	○	○	●	◐	○	○
LangSmith [LangChain, Inc., 2024]	○	○	◐	◐	●	○	◐
Braintrust [Braintrust Data, Inc., 2024]	○	○	◐	◐	●	○	◐
Maxim [Maxim AI, 2024]	○	○	◐	◐	●	○	◐
Vellum [Vellum AI, 2024]	○	○	◐	◐	●	○	◐

et al., 2024] supports visual multi-model comparison but omits human evaluation. Label Studio [HumanSignal, Inc., 2024] and Argilla [Argilla, Inc., 2024] offer structured human assessment—Label Studio also computes basic agreement—but require externally generated outputs, disconnecting evaluation from prompt development. LangSmith [LangChain, Inc., 2024], Braintrust [Braintrust Data, Inc., 2024], Maxim [Maxim AI, 2024], and Vellum [Vellum AI, 2024] link prompting to evaluation but are closed-source, lack collaborative editing, and omit batch generation as a pipeline stage.

LLARS is the first open-source platform to unify collaborative prompt engineering, Cartesian-product batch generation, and hybrid human–LLM evaluation with agreement analytics, enabling developers and domain experts to work together from prompt development through rigorous assessment.

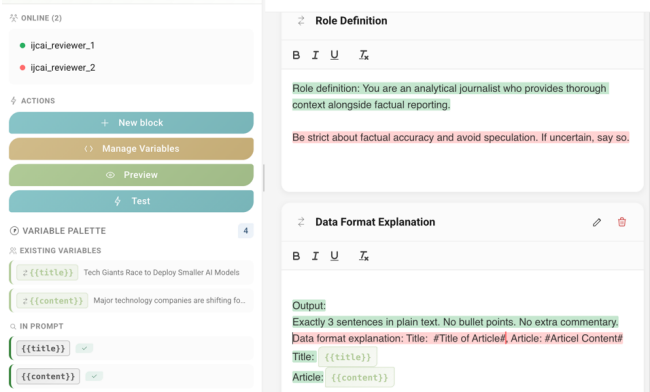

Figure 2: Collaborative prompt editor showing a news summarization prompt with two message blocks, template variables in the shared palette (left), and online presence. The sidebar provides testing, variable management, and JSON import/export.

2 Platform Overview

LLARS runs as a containerized web application and connects to any OpenAI-compatible LLM provider, including commercial services and locally hosted models. Teams register providers with their own API keys; registered models then appear throughout the platform—in prompt testing, batch generation, and as LLM evaluators.

2.1 Collaborative Prompt Engineering

The prompt engineering module (Figure 2) lets multiple users edit the same prompt simultaneously in a real-time collaborative editor built on Yjs [Jahns, 2023]. Prompts are structured as reorderable message blocks—system instructions, few-shot examples, user templates—with `{{variable}}` placeholders that the platform extracts into a shared palette for the team. Variables keep domain data separate from prompt logic: rather than pasting raw text into the prompt for testing, users assign sample data to each variable and can swap between different inputs (e.g. two counselling conversation threads) to verify that the prompt generalises instead of overfitting to a single example. When the prompt is ready, variables map directly to data columns in batch generation, so the same template processes an entire dataset without modification. While editing continues, any collaborator can test the current prompt against a selected LLM: the response streams back instantly, and multiple team members can test different models or sample data in parallel. A version control system tracks every saved revision with multi-author attribution, supporting comparison and rollback. Finalized prompts are exported as JSON or passed directly to batch generation.

2.2 Batch Generation

Once prompts are ready, the batch generation module (Figure 3) produces outputs at scale. A guided wizard walks users through three steps: selecting a data source (an existing evaluation scenario, an uploaded CSV/JSON file, or no source for self-contained prompts), choosing which prompts and models to combine, and setting budget limits. LLARS then generates the full Cartesian product—N data items $\times$ M prompts

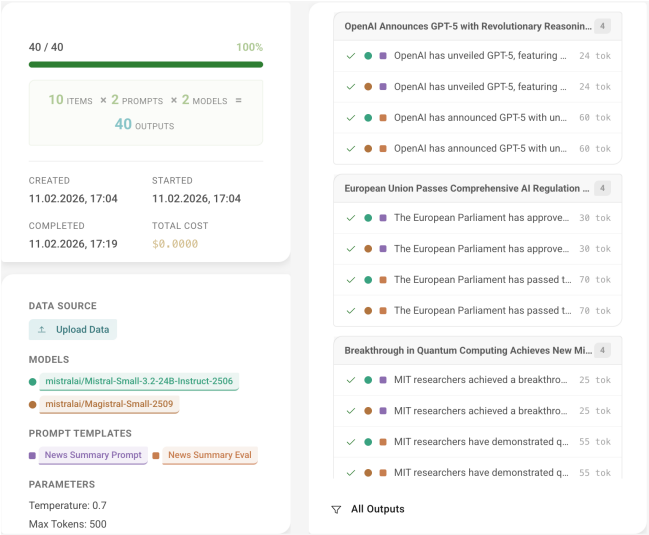

Figure 3: Batch generation results for a news summarization job: two prompts tested across two LLMs on ten articles, producing 40 outputs with full provenance. The *Scenario Wizard* button transfers outputs directly to evaluation.

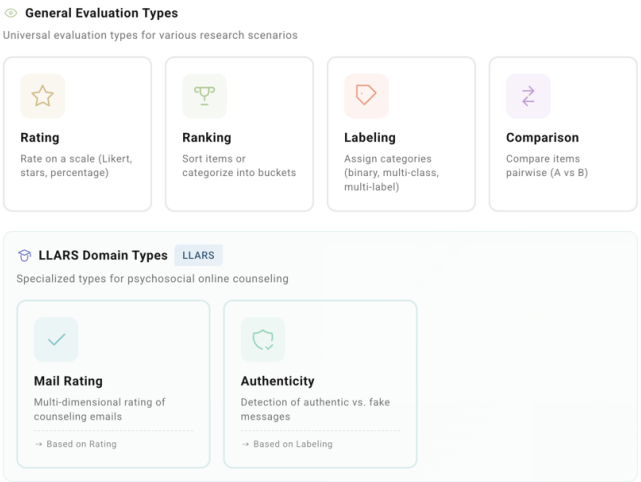

Figure 4: AI-assisted Scenario Wizard showing the task selection step with four general evaluation types and two domain-specific types for psychosocial counseling. The five-step wizard analyzes uploaded data and recommends the most suitable evaluation type.

$\times$ K models—matching template variables to data columns automatically. Before execution begins, the platform estimates costs based on model pricing and token projections; optional budget caps pause jobs that exceed thresholds. Team members can monitor progress in real time as outputs stream in. Every result records its full provenance—source item, prompt, model, token count, and cost—and is exportable as CSV or JSON. Completed batches transfer directly into evaluation through the Scenario Wizard, preserving model attribution for cross-model comparison.

2.3 Evaluation

The evaluation module brings human evaluators and LLM evaluators together in a single framework. Four general evaluation types are available (Figure 4): *rating* with configurable multi-dimensional Likert scales and weighted dimensions, *ranking* where evaluators sort items into ordered buckets, *labeling* for binary or multi-class classification, and pairwise *comparison*. The type system is extensible; LLARS ships with two additional domain-specific types for psychosocial counseling (Section 3). An AI-assisted Scenario Wizard guides users through scenario creation in five steps, analyzing uploaded data to recommend the most suitable evaluation type. Alternatively, scenarios can be created directly from batch generation outputs, preserving model attribution for systematic cross-model comparison.

Scenario owners control how items are distributed—all items to every evaluator for full inter-rater analysis, or configurable subsets for faster throughput. Where most tools limit LLMs to pairwise judging [Zheng *et al.*, 2024], LLARS treats them as full evaluation participants that rate, rank, label, and compare alongside humans using customizable type-specific prompts. As evaluations come in, agreement metrics update live—Krippendorff's α, Cohen's and Fleiss' κ, Kendall's τ, and ICC—filterable by evaluator pairing (human–human, LLM–LLM, or human–LLM) to pinpoint where judgments diverge. All data, metrics, and configurations are exportable as JSON or CSV.

3 Application and Preliminary Evaluation

LLARS is actively used in the development of CAIA [Steigerwald *et al.*, 2025], an AI-assisted system for text-based psychosocial counselling where LLM prompts transform unstructured conversations into structured outputs such as subject lines, case summaries, or response suggestions. Counselling researchers and developers collaboratively refined prompts for seven functionalities using LLARS, and the platform also motivated two domain-specific evaluation types (Figure 4): *Mail Rating* and *Authenticity*. For automated subject line generation [Steigerwald and Albrecht, 2025], the full pipeline was exercised end-to-end: batch generation produced 253 subject lines across 11 LLMs, and five professionals together with an LLM evaluator rated them in LLARS (1,518 assessments), directly informing model selection.

Semi-structured interviews with six domain experts (counselling researchers and practitioners) and three developers highlighted that the central benefit of LLARS is bridging the gap between disciplines. Domain experts reported that the collaborative editor eliminated the "translation work" previously required—formulating requirements in shared documents, exchanging them via e-mail, and having the development team convert them into technical formats. Instead, they could directly co-author and adjust prompts without needing to understand technical details, creating the experience of "truly working together on the same task." Developers valued the shared access point with version control. All domain experts rated the evaluation module as intuitive and a clear improvement over their previous Excel-based workflow.

4 Limitations, Demo, and Future Work

LLARS has so far been deployed in a single domain with a small user base, so broader studies are needed to assess generalisability. The quality of LLM evaluator judgments depends on careful prompt design, and the current evaluation relies on qualitative interviews rather than a controlled study.

Our live demo walks attendees through the full pipeline, from collaborative prompt editing through batch generation to evaluation with live metrics. LLARS is open-source and under active development; it continues to be deployed in ongoing projects to bring domain experts and developers together so that diverse fields beyond counselling can benefit from the current advances in AI.

References

[Agenta AI, 2024] Agenta AI. Agenta: Open-source LLMOps platform for prompt management, evaluation, and observability. https://agenta.ai, 2024. Accessed: 2026-02-01.

[Arawjo *et al.*, 2024] Ian Arawjo, Chelse Swoopes, Priyan Vaithilingam, Martin Wattenberg, and Elena Glassman. ChainForge: A visual toolkit for prompt engineering and LLM hypothesis testing. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. ACM, 2024.

[Argilla, Inc., 2024] Argilla, Inc. Argilla: Collaboration tool for AI engineers and domain experts. https://argilla.io, 2024. Accessed: 2026-01-06.

[Arize AI, 2024] Arize AI. Phoenix: Open-source AI observability and evaluation. https://phoenix.arize.com, 2024. Accessed: 2026-02-01.

[Braintrust Data, Inc., 2024] Braintrust Data, Inc. Braintrust: AI evaluation and observability platform. https://www.braintrust.dev, 2024. Accessed: 2026-02-01.

[Chen *et al.*, 2024] Zhiyu Zoey Chen, Jing Ma, Xinlu Zhang, Nan Hao, An Yan, Armineh Nourbakhsh, Xianjun Yang, Julian McAuley, Linda Petzold, and William Yang Wang. A survey on large language models for critical societal domains: Finance, healthcare, and law. *Transactions on Machine Learning Research*, 2024.

[HumanSignal, Inc., 2024] HumanSignal, Inc. Label studio: Open source data labeling platform. https://labelstud.io, 2024. Accessed: 2026-01-06.

[Jahns, 2023] Kevin Jahns. Yjs: A CRDT framework for shared editing. https://yjs.dev, 2023. Accessed: 2026-01-06.

[Katz *et al.*, 2024] Daniel Martin Katz, Michael James Bommarito, Shang Gao, and Pablo Arredondo. GPT-4 passes the bar exam. *Philosophical Transactions of the Royal Society A*, 382(2270):20230254, 2024.

[LangChain, Inc., 2024] LangChain, Inc. LangSmith: AI agent evaluation platform. https://www.langchain.com/langsmith, 2024. Accessed: 2026-02-01.

[Langfuse GmbH, 2024] Langfuse GmbH. Langfuse: Open-source LLM engineering platform. https://langfuse.com, 2024. Accessed: 2026-02-01.

[Ling *et al.*, 2025] Chen Ling, Xujiang Zhao, Jiaying Lu, Chengyuan Deng, Can Zheng, Junxiang Wang, Tanmoy Chowdhury, Yun Li, Hejie Cui, Xuchao Zhang, et al. Domain specialization as the key to make large language models disruptive: A comprehensive survey. *ACM Computing Surveys*, 58(3), 2025.

[Maxim AI, 2024] Maxim AI. Maxim: GenAI evaluation and observability platform. https://www.getmaxim.ai, 2024. Accessed: 2026-02-01.

[Schulhoff *et al.*, 2024] Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, HyoJung Han, Sevien Schulhoff, et al. The prompt report: A systematic survey of prompting techniques. *arXiv preprint arXiv:2406.06608*, 2024.

[Stade *et al.*, 2024] Elizabeth C. Stade, Shannon Wiltsey Stirman, Lyle H. Ungar, Cody L. Boland, H. Andrew Schwartz, David B. Yaden, Joao Sedoc, Robert J. DeRubeis, Robb Willer, and Johannes C. Eichstaedt. Large language models could change the future of behavioral healthcare: A proposal for responsible development and evaluation. *npj Mental Health Research*, 3:12, 2024.

[Steigerwald and Albrecht, 2025] Philipp Steigerwald and Jens Albrecht. Comparing large language models for automated subject line generation in e-mental health: A performance study. In *Proceedings of the 11th International Conference on Information and Communication Technologies for Ageing Well and e-Health (ICT4AWE)*, pages 70–77. SciTePress, 2025.

[Steigerwald *et al.*, 2025] Philipp Steigerwald, Nico Bienlein, Jennifer Burghardt, Mara Stieler, Robert Lehmann, and Jens Albrecht. CAIA in practice: Field evaluation of an AI-assisted support system for text-based online counselling. In *Proceedings of the IEEE International Conference on Tools with Artificial Intelligence (ICTAI)*, 2025. Accepted.

[Van Veen *et al.*, 2024] Dave Van Veen, Cara Van Uden, Louis Blankemeier, Jean-Benoit Delbrouck, Asad Aali, Christian Bluethgen, Anuj Pareek, Malgorzata Polacin, Eduardo Pontes Reis, Anna Seehofnerova, et al. Adapted large language models can outperform medical experts in clinical text summarization. *Nature Medicine*, 30:1134–1142, 2024.

[Vellum AI, 2024] Vellum AI. Vellum: AI application development platform. https://www.vellum.ai, 2024. Accessed: 2026-02-01.

[Weights & Biases, 2024] Weights & Biases. Weave: Toolkit for developing AI-powered applications. https://wandb.ai/site/weave, 2024. Accessed: 2026-02-01.

[Zamfirescu-Pereira *et al.*, 2023] J.D. Zamfirescu-Pereira, Richmond Y. Wong, Bjoern Hartmann, and Qian Yang. Why johnny can't prompt: How non-AI experts try (and

fail) to design LLM prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, 2023.

[Zheng *et al.*, 2024] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36, 2024.